

## Exercise 2.6: Create and customize a systemd service

This exercise is going to explore the various configuration directories for a **systemd** service. The application **stress** will be used, install the package **stress** with the appropriate package installer and examine the files installed. Notice the absolute path name of the binary **stress** and the absence of a **stress.service** file in the **systemd** file structure.



### Very Important

Some distributions have the **stress** package; others have the newer **stress-ng**. Either package will work as **stress-ng** supports the same command line options as **stress**.



### On CentOS

```
# yum install stress
# rpm -ql stress
```



### On openSUSE

```
# zypper install stress-ng
# rpm -ql stress-ng
```



### On Ubuntu

```
# apt-get install stress
# dpkg -L stress
```

The **stress** package does not include a **systemd** unit configuration, so one must be created. The package installed on the test system has the binary for **stress** as `/usr/bin/stress`. Create a **systemd** vendor unit file as `/usr/lib/systemd/system/foo.service`. You will require **root** level access to create a file in this directory.

```
$ sudo vim /usr/lib/systemd/system/foo.service
```



### `/usr/lib/systemd/system/foo.service`

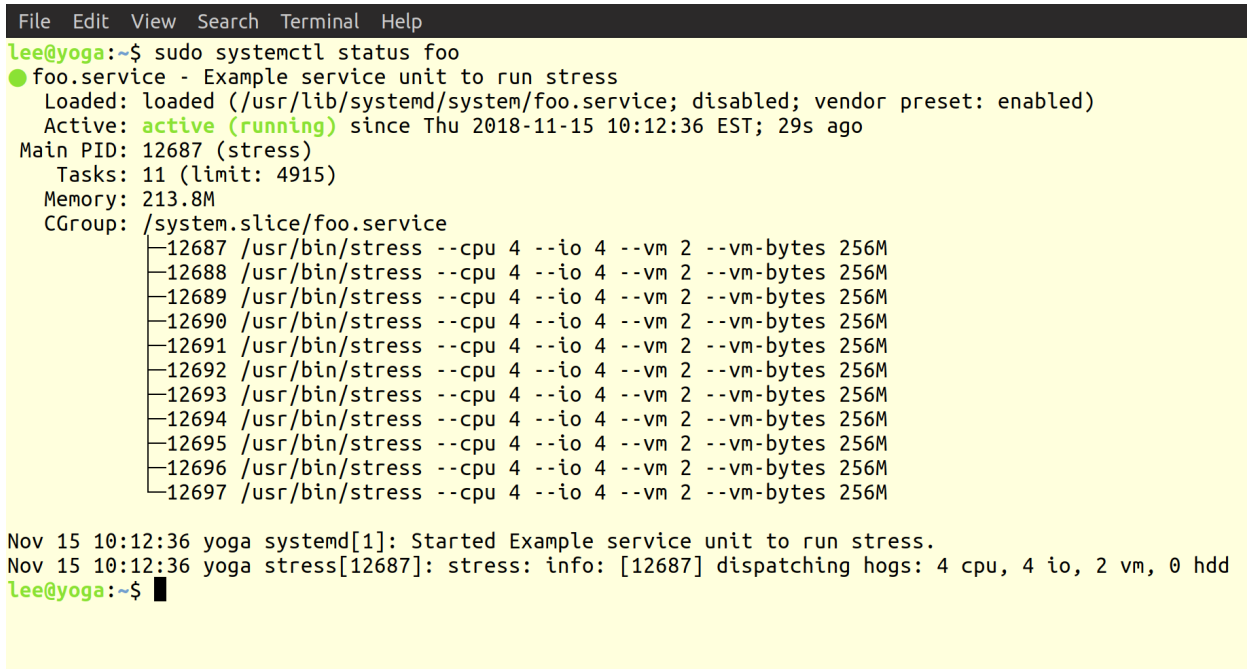
```
[Unit]
Description=Example service unit to run stress
[Service]
ExecStart=/usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
[Install]
WantedBy=multi-user.target
```

A copy of this file (`foo1.service`) can be found in the tarball in the `LFS211/SOLUTIONS/s_02/` directory.

Once the unit file is created **systemd** will be able to start and stop the service. Use the **top** command to verify that **stress** is working. The following commands may be useful:

```
# systemctl daemon-reload
# systemctl start foo
# systemctl status foo -l
# systemd-delta
# systemctl stop foo
```

The example program **stress** which is now a service, does not display much feedback as to what it is doing. The **systemctl status** of the service can be checked, the output would look something like this:



```
File Edit View Search Terminal Help
lee@yoga:~$ sudo systemctl status foo
● foo.service - Example service unit to run stress
   Loaded: loaded (/usr/lib/systemd/system/foo.service; disabled; vendor preset: enabled)
   Active: active (running) since Thu 2018-11-15 10:12:36 EST; 29s ago
 Main PID: 12687 (stress)
    Tasks: 11 (limit: 4915)
   Memory: 213.8M
    CGroup: /system.slice/foo.service
            └─12687 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
              └─12688 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                └─12689 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                  └─12690 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                    └─12691 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                      └─12692 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                        └─12693 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                          └─12694 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                            └─12695 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                              └─12696 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
                                └─12697 /usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M

Nov 15 10:12:36 yoga systemd[1]: Started Example service unit to run stress.
Nov 15 10:12:36 yoga stress[12687]: stress: info: [12687] dispatching hogs: 4 cpu, 4 io, 2 vm, 0 hdd
lee@yoga:~$
```

Figure 2.18: **systemctl status foo**

Examining the output of the **top** command will show two processes of **stress** with more than the specified 256M of memory, four processes of **stress** with nearly 100% CPU usage and four processes with neither high CPU or high memory usage. They would be the memory hogs, CPU hogs, and io hogs, respectively. See the example below:

File Edit View Search Terminal Help										
top - 10:14:39 up 2:39, 1 user, load average: 9.14, 4.06, 2.87										
Tasks: 303 total, 10 running, 293 sleeping, 0 stopped, 0 zombie										
%Cpu(s): 56.0 us, 43.3 sy, 0.0 ni, 0.0 id, 0.6 wa, 0.0 hi, 0.0 si, 0.0 st										
MiB Mem : 15786.5 total, 8109.7 free, 3808.1 used, 3868.7 buff/cache										
MiB Swap: 2048.0 total, 2048.0 free, 0.0 used. 10771.2 avail Mem										
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+ COMMAND
12691	root	20	0	4056	100	0	R	99.0	0.0	1:59.64 stress
12694	root	20	0	4056	100	0	R	98.7	0.0	2:01.10 stress
12688	root	20	0	4056	100	0	R	98.0	0.0	1:59.32 stress
12696	root	20	0	4056	100	0	R	97.4	0.0	1:59.99 stress
12693	root	20	0	266204	133160	272	R	96.0	0.8	1:59.42 stress
12690	root	20	0	266204	90656	272	R	94.4	0.6	1:59.67 stress
12695	root	20	0	4056	100	0	D	44.0	0.0	0:44.99 stress
12689	root	20	0	4056	100	0	R	43.4	0.0	0:43.68 stress
12692	root	20	0	4056	100	0	R	42.1	0.0	0:43.82 stress
12697	root	20	0	4056	100	0	R	40.7	0.0	0:43.50 stress
2564	lee	20	0	3767724	424628	158296	S	16.9	2.6	6:57.38 gnome-she+
447	root	0	-20	0	0	0	I	4.6	0.0	0:12.16 kworker/1+
2396	lee	20	0	866864	208236	154892	S	4.3	1.3	4:39.05 Xorg
233	root	-51	0	0	0	0	S	2.3	0.0	0:46.09 irq/51-SY+
300	root	0	-20	0	0	0	I	2.3	0.0	0:04.81 kworker/4+
13093	lee	20	0	508924	36068	28092	S	1.7	0.2	0:00.55 gnome-scr+
68	root	25	5	0	0	0	S	1.0	0.0	0:37.74 ksmd

Figure 2.19: top

As we are interested specifically in the **stress** service and its child processes, we provide a script for monitoring the service processes running.

**track-STRESS.sh** looks for a **stress** process with a PPID of 1, then all of the related child processes. Once the processes are located some data is extracted with the **ps** command.

A copy of this script ([track-STRESS.sh](#)) can be found in the tarball in the [LFS211/SOLUTIONS/s\\_02/](#) directory.

SH

**track-STRESS.sh**

```
#!/bin/bash
#
# This little script is used with the LFS311 lab exercise
# on systemd startup files and their affect on a background
# service. The "foo" service.
#
# The script looks for "stress" launched with a PPID of 1
# and it's children processes then uses the "ps" command
# to collect some information.

while [ true ]
do

# reset the variables to make sure we are picking up
# current information
PID=""
PID1=""
clear

echo "$0 is running"

# sift and sort the pid's
PID1=` ps -ef | grep stress | grep -v grep | awk '{print $1,$2,$3,$6}' `
PID=`echo $PID1 | grep "1 "| awk '{print $2}' `
echo "The pid for stress is $PID"
```

SH

```
# using our list of pid's, grab some information
ps --ppid $PID --pid $PID -o pid,ppid,comm,vsz,pcpu,psr,slice 2>/dev/null

sleep 5
done
exit
```

An example of the script running:

```
File Edit View Search Terminal Tabs Help
lee@yoga: ~/LFS/LF-Current-LFS3... x root@yoga: /home/lee/LFS/LF-C... x
./track-STRESS.sh is running
The pid for stress is 12687
  PID  PPID  COMMAND      VSZ  %CPU  PSR  SLICE
12687    1  stress      4056   0.0   7  system.slice
12688 12687  stress      4056  97.6   2  system.slice
12689 12687  stress      4056  36.0   5  system.slice
12690 12687  stress     266204  97.4   7  system.slice
12691 12687  stress      4056  97.7   1  system.slice
12692 12687  stress      4056  36.6   2  system.slice
12693 12687  stress     266204  97.4   6  system.slice
12694 12687  stress      4056  98.9   3  system.slice
12695 12687  stress      4056  37.8   5  system.slice
12696 12687  stress      4056  97.6   4  system.slice
12697 12687  stress      4056  36.5   5  system.slice
```

Figure 2.20: **track-STRESS.sh**

Since `/usr/lib/systemd/system/foo.service` is the default configuration supplied by the packager of the service and may be altered by the vendor at any time, create a custom unit file in `/etc/systemd/system/foo.service` for the **stress** service. This file is not usually overwritten by the vendor so local customizations can go here. Change the parameters slightly for the **foo** service using this directory. It is common practice to copy the vendor unit file into the `/etc/systemd/system/` directory and make appropriate customizations.



### `/etc/systemd/system/foo.service`

```
[Unit]
Description=Example service unit to run stress
[Service]
ExecStart=/usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
[Install]
WantedBy=multi-user.target
```

A copy of this file (`foo2.service`) can be found in the tarball in the `LFS211/SOLUTIONS/s_02/` directory.

Start or restart the service and examine the differences in the following commands output.

```
# track-STRESS.sh
# systemctl status foo -l
# systemd-delta
```

The changes to the configuration file can be seen with the **track-STRESS.sh** script, notice the number of memory hogs is

now 4 and the CPU hogs is reduced to 2.

```
File Edit View Search Terminal Help
./track-STRESS.sh is running
The pid for stress is 17860
  PID PPID COMMAND      VSZ %CPU PSR SLICE
17860   1 stress          4056 0.0  4 system.slice
17861 17860 stress          4056 99.4  0 system.slice
17862 17860 stress          4056 43.7  1 system.slice
17863 17860 stress        266204 99.3  2 system.slice
17864 17860 stress          4056 99.5  5 system.slice
17865 17860 stress          4056 43.9  1 system.slice
17866 17860 stress        266204 99.4  4 system.slice
17867 17860 stress        266204 99.3  3 system.slice
17868 17860 stress        266204 99.3  6 system.slice
```

Figure 2.21: **track-STRESS.sh** with new unit file

Which configuration (or **unit** file) file is active is not clear in the script. Use `systemctl status foo` to see which unit file is being used. This will show which configuration files are being used but not the differences in the files.

```
File Edit View Search Terminal Help
lee@yoga:~$ sudo systemctl status foo
● foo.service - Example service unit to run stress
   Loaded: loaded (/etc/systemd/system/foo.service; disabled; vendor preset: enabled)
   Active: active (running) since Thu 2018-11-15 11:36:45 EST; 1min 17s ago
 Main PID: 17860 (stress)
    Tasks: 9 (limit: 4915)
   Memory: 482.7M
    CGroup: /system.slice/foo.service
            └─17860 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
              └─17861 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                └─17862 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                  └─17863 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                    └─17864 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                      └─17865 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                        └─17866 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                          └─17867 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
                            └─17868 /usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M

Nov 15 11:36:45 yoga systemd[1]: Started Example service unit to run stress.
Nov 15 11:36:45 yoga stress[17860]: stress: info: [17860] dispatching hogs: 2 cpu, 2 io, 4 vm, 0 hdd
lee@yoga:~$
```

Figure 2.22: **systemctl status foo** changing the unit file

To see the details of the unit file changes the **systemd-delta** command can be used. The output has the changes in **diff** format, making it easy to see what has changed. See the example below:

```

File Edit View Search Terminal Help

[OVERRIDDEN] /etc/systemd/system/foo.service → /usr/lib/systemd/system/foo.service

--- /usr/lib/systemd/system/foo.service 2018-11-15 09:24:28.183557249 -0500
+++ /etc/systemd/system/foo.service      2018-11-15 11:34:21.689073568 -0500
@@ -1,6 +1,6 @@
[Unit]
Description=Example service unit to run stress
[Service]
-ExecStart=/usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
+ExecStart=/usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
[Install]
WantedBy=multi-user.target

[EXTENDED] /lib/systemd/system/rc-local.service → /lib/systemd/system/rc-local.service.d/debian.con
[EXTENDED] /lib/systemd/system/user@.service → /lib/systemd/system/user@.service.d/timeout.conf

4 overridden configuration files found.
lines 8-25/25 (END)

```

Figure 2.23: **systemd-delta** showing unit file override

Often times it is desirable to add or change features by program or script control, the drop-in files are convenient for this. One item of caution, if one is changing a previously defined function (like `ExecStart`) it must be undefined first then added back in. Create a drop-in directory and file for out **stress** service and verify the changes are active. Our example file for `foo.service` using a drop-in directory (`00-foo.conf`) can be found in the SOLUTIONS tarball in the `LFS211/SOLUTIONS/s_02/` directory and contains:



```
/etc/systemd/system/foo.service.d/00-foo.conf
```

```

[Service]
ExecStart=
ExecStart=/usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M

```

Start or restart the service and examine the differences in the output of the following commands.

```

# track-STRESS.sh
# systemctl status foo -l
# systemd-delta

```

The information in the drop in file over writes the unit file. In this example the number of “hogs” has been greatly reduced.

```

File Edit View Search Terminal Help

./track-STRESS.sh is running
The pid for stress is 8044
  PID  PPID  COMMAND      VSZ  %CPU  PSR  SLICE
  8044    1  stress       4056  0.0   3  system.slice
  8045  8044  stress       4056  99.9   1  system.slice
  8046  8044  stress       4056  31.4   6  system.slice
  8047  8044  stress     135132  99.8   7  system.slice

```

Figure 2.24: **track-STRESS.sh** using dropin file

**systemctl status** shows the dropin file is active. If there was several dropin files it would show the order that were applied to

the service.

```
File Edit View Search Terminal Help
root@yoga:/etc/systemd/system/foo.service.d# systemctl status foo
● foo.service - Example service unit to run stress
   Loaded: loaded (/etc/systemd/system/foo.service; disabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/foo.service.d
            └─00-foo.conf
   Active: active (running) since Thu 2018-11-15 14:06:05 EST; 13s ago
   Main PID: 8044 (stress)
     Tasks: 4 (limit: 4915)
    Memory: 103.5M
    CGroup: /system.slice/foo.service
            └─8044 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M
              └─8045 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M
                └─8046 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M
                  └─8047 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M

Nov 15 14:06:05 yoga systemd[1]: Started Example service unit to run stress.
Nov 15 14:06:05 yoga stress[8044]: stress: info: [8044] dispatching hogs: 1 cpu, 1 io, 1 vm, 0 hdd
root@yoga:/etc/systemd/system/foo.service.d#
```

Figure 2.25: **systemctl status** showing unit file override and dropin file

Like the other commands, **systemd-delta** shows the files used by the service. In addition to the files used, the details of the changes in the unit file are displayed. Notice the changes to the service made with the drop-in file are not displayed, only the file name.

```
File Edit View Search Terminal Help

[OVERRIDDEN] /etc/systemd/system/foo.service → /usr/lib/systemd/system/foo.service

--- /usr/lib/systemd/system/foo.service 2018-11-15 09:24:28.183557249 -0500
+++ /etc/systemd/system/foo.service      2018-11-15 11:34:21.689073568 -0500
@@ -1,6 +1,6 @@
[Unit]
Description=Example service unit to run stress
[Service]
-ExecStart=/usr/bin/stress --cpu 4 --io 4 --vm 2 --vm-bytes 256M
+ExecStart=/usr/bin/stress --cpu 2 --io 2 --vm 4 --vm-bytes 256M
[Install]
WantedBy=multi-user.target

[EXTENDED] /lib/systemd/system/rc-local.service → /lib/systemd/system/rc-local.service.d/debian.conf
[EXTENDED] /lib/systemd/system/user@.service → /lib/systemd/system/user@.service.d/timeout.conf

4 overridden configuration files found.
lines 8-25/25 (END)
```

Figure 2.26: **systemd-delta** showing unit file and dropin file

With **systemd**, additional features and capabilities can be easily added. As an example, **cgroups** controls can be added to our service. Here is an example of adding a **systemd slice** to the example service and adding a resource limit to that slice. The slice is then attached to the service drop-in file. First setup a `<service>.slice` unit file:



**/etc/systemd/system/foo.slice**

```
[Unit]
Description=stress slice
[Slice]
CPUQuota=30%
```



A copy of this file (`foo.slice`) can be found in the tarball in the `LFS211/SOLUTIONS/s_02/` directory.

Then connect our service to the **slice**. Add the following to the bottom of the unit file in `/etc/systemd/system/foo.service.d/00-foo.conf`



### connect slice file to service file

```
Slice=foo.slice
```

Restart the services and examine the differences with:

```
# systemctl stop foo
# systemctl daemon-reload
# systemctl start foo
# systemctl status foo -l
# systemd-delta
# top
# track-STRESS.sh
```

The cgroup information in the **slice** has been applied to the service. Notice the amount of CPU resource consumed. The total is 30% of one processor but it may be spread across multiple CPU's.

File Edit View Search Terminal Help						
./track-STRESS.sh is running						
The pid for stress is 8989						
PID	PPID	COMMAND	VSZ	%CPU	PSR	SLICE
8989	1	stress	4056	0.0	7	foo.slice
8990	8989	stress	4056	12.7	5	foo.slice
8991	8989	stress	4056	5.0	4	foo.slice
8992	8989	stress	135132	12.0	3	foo.slice

Figure 2.27: track-STRESS.sh using slice attribute

The **systemctl status** command shows the change in **CGroup** with the **slice** attribute active.

```
File Edit View Search Terminal Help
root@yoga:/etc/systemd/system# systemctl status foo
● foo.service - Example service unit to run stress
   Loaded: loaded (/etc/systemd/system/foo.service; disabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/foo.service.d
            └─00-foo.conf
   Active: active (running) since Thu 2018-11-15 14:30:29 EST; 2min 53s ago
 Main PID: 8989 (stress)
   Tasks: 4 (limit: 4915)
  Memory: 10.4M
   CGroup: /foo.slice/foo.service
            └─8989 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M
              └─8990 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M
                └─8991 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M
                  └─8992 /usr/bin/stress --cpu 1 --vm 1 --io 1 --vm-bytes 128M

Nov 15 14:30:29 yoga systemd[1]: Started Example service unit to run stress.
Nov 15 14:30:29 yoga stress[8989]: stress: info: [8989] dispatching hogs: 1 cpu, 1 io, 1 vm, 0 hdd
root@yoga:/etc/systemd/system#
```

Figure 2.28: systemctl status showing CGroup and slice

**Bonus step:** In our example there are no unique values in the `/etc/systemd/system/foo.service` file so in this example



it is redundant. We can get rid of the extra file.

```
# mv /etc/systemd/system/foo.service /root/  
# systemctl daemon-reload  
# systemctl restart foo  
# systemctl status foo
```

Consult the man pages **systemd.resource-control(5)**, **systemd.service(5)**, **systemd-delta(1)** and other **systemd** man pages for additional information.